



hmn
ai masterclass
for decisionmakers

hmn



Alex Schwartz

CEO, hmn

- Former Publicis Sapient strategy director; led enterprise transformation for Fortune 500s
- Founded The Disrupted Workforce & built hmn
- Advises executives & teams on closing gap between AI ambition and organizational readiness



Matty Squarzoni

CTO, hmn

- Architect of hmn's AI assessment platform & diagnostic tooling
- Builds systems for real-time visibility into workforce AI readiness
- Skills benchmarking, adoption tracking, scalable technology



Dr. Ernesto Lee

Chief AI Strategist, hmn

- #1 AI Professor in America (2025); founder of DrLee.AI
- Advises C-suite at JP Morgan Chase, Cisco, and federal agencies
- Designs assessment frameworks bridging AI literacy & enterprise adoption



Human-AI Synergy Is The Future

AI is most powerful when it amplifies human capability. Organizations mastering this synergy will crush those who debate it.



Culture Eats Strategy for Breakfast

Your million-dollar AI strategy means nothing if your culture kills it. Transformation lives or dies at the water cooler and on Slack, not the boardroom.



Adaptability is the Only Constant

In a world of exponential change, the ability to continuously learn, and unlearn, matters more than skills or degrees.



Your Tech Stack is Worthless Unless People Get It

Technology is only as powerful as the humans who wield it. We bet on people first. Everything else follows.



The fastest rate of change in human history

“

20 Years from now, the rate
rate of change will be 4x what
what it is currently

— Ray Kurzweil



“In 2022, AI couldn't do basic arithmetic reliably. It would confidently tell you that $7 \times 8 = 54$.

By 2023, it could pass the bar exam.

By 2024, it could write working software and explain graduate-level science.

By late 2025, some of the best engineers in the world said they had handed over most of their coding work to AI.

On February 5th, 2026, new models arrived that made everything before them feel like a different era.”

— Matt Shumer, “Something Big Is Happening” 2/9/26

75%



AI capabilities now double every 6 months. **75% of employees fear their skills will be obsolete in 3–5 years.**

The Gap Is Real. The Opportunity Is Now.

You don't need to be technical. You need to start.

"We're not experiencing an AI failure. We're experiencing a measurement failure." — Gallacher, UC Berkeley BRIE, 2025

1.7x

revenue growth for AI leaders
vs. those who delay

BCG, 2025

Only 5%

of companies are "future-built" for AI.
60% report minimal or no value.

BCG, 2025

55%

of AI winners redesigned workflows.

McKinsey, 2025

15%

of workers over 45 use AI tools
— but 50% want training

WorkLife / Randstad, 2025

50+

internal apps built by non-technical
users at Sears — no engineers

a16z, 2025

Days → Seconds

investment research compressed
with AI-powered workflows

LangChain / Captide, 2025

Before You Drive the Ferrari

Fasten your seatbelt first.



Photo: Unsplash

- **AI is confident when it's wrong.** It doesn't hesitate. It doesn't say "I'm guessing." Verify everything.
- **Garbage in, garbage out.** Vague prompts get vague answers. Specificity is a skill.
- **Don't feed it your secrets.** No client data, no financials, no confidential info — yet.
- **It doesn't know what you don't tell it.** AI has no context about your business until you give it context.
- **Digital hygiene comes first.** If your files, contacts, and data are a mess, AI amplifies the mess.
- **Known problems are real.** Bias, hallucination, and privacy gaps are features of the current moment.
- **Start with one workflow. Master it. Then expand.** The biggest mistake is trying to do everything at once.
- **NEVER let AI think for you.** This is one of the greatest pitfalls. Your judgment is the product, not the output.
- **We're early. The tools change fast.** Learn the thinking, not just the tool.

1. Role Based Prompting (Be The Expert I Need)

"Act as a cynical, veteran Chief Marketing Officer with 20 years of crisis management experience in the tech sector. Review this PR strategy."

Reality Check: By default, an LLM acts as a generalized assistant. Left undirected, it lacks the specialized lens required for high-level consulting.

2. Chain of Thought Reasoning (Tell Me Why)

"Before you give me the final answer, walk me through your reasoning step-by-step." Today is April 3rd, 2026

Reality Check: AI naturally jumps to the most obvious, statistically likely conclusion. If you don't force it to "think out loud," you risk getting shallow, generic answers.

3. Few Shot Prompting (Giving Examples)

"Here are my best and worst examples..."

Reality Check: "Zero-shot" prompting—asking a cold question with no examples—yields average, uninspired results.

4. Reverse Prompting (Engineering Outcomes)

"Ask me any questions you need to perfectly understand my request."

Reality Check: We are often blind to what context we leave out of a brief. Assuming you've given the AI everything it needs is a massive vulnerability.

5. Sparring Partner (Don't Flatter Me)

"I am going to pitch a \$500k digital transformation project. You play the role of a highly skeptical, budget-conscious CFO. Push back on my ideas one by one and don't flatter me. Let's begin."


The Reality Check: Treating AI purely as a static text generator ignores its most powerful capability: dynamic, real-time simulation.

THE TOKEN:

The basic unit of AI understanding.



2M Advertised Limit
(actual)



EQUAL TO
~750,000
WORDS




RULE OF THUMB:
1 Token $\approx$
0.75 Words

ESTABLISHING SCALE:
1 MILLION
TOKENS

BUSINESS CONTEXT:
ANALYZING ~15-20
STANDARD COMMERCIAL
LEASE AGREEMENTS

THE "LOST IN THE MIDDLE" PROBLEM



STRATEGIC CRE SOLUTION: TARGETED RETRIEVAL (RAG)

DUMP ALL DATA APPROACH



STRATEGIC "RAG" APPROACH



FRONTIER - HIGH REASONING:
GPT-5.2 Pro

Context Window: 400,000 tokens
(Approx. 300,000 words)



INPUT COST
(per 1M tokens)



OUTPUT COST
(per 1M tokens)



IDEAL CRE USE CASE:



Complex Contract
Review &
Portfolio Strategy

FRONTIER - MULTIMODAL EFFICIENCY:
Gemini 3.1 Pro

Context Window: 2 Million tokens
(Approx. 1.5M words)



INPUT COST
(per 1M tokens)



OUTPUT COST
(per 1M tokens)



IDEAL CRE USE CASE:



Multimodal Analysis
(Site Photos & Cross-Referencing Maps)

FRONTIER - NUANCED ANALYTICS:
Claude 4.6 Opus

Context Window: 1 Million tokens
(Approx. 750,000 words)



INPUT COST
(per 1M tokens)



OUTPUT COST
(per 1M tokens)



IDEAL CRE USE CASE:



Complex Text
Synthesis
(Regulations,
Reporting)



Long-Form
Reporting
(Feasibility,
Market Trends)

STRATEGY: Optimize Costs by Aligning Task Complexity with the Appropriate Model Tier.

* Upgrade for Pro HD